


# Packet Sniffing and Spoofing

Ethical Hacking Challenge #1

**Francesco Marchiori**

ID = 2020389



A.Y. 2021/2022

## Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
<b>2</b>	<b>Environment Setup</b>	<b>1</b>
2.1	Network Interface Name . . . . .	1
<b>3</b>	<b>Task 1</b>	<b>1</b>
3.1	Task 1.1: Using Scapy to Sniff and Spoof Packets . . . . .	2
3.2	Task 1.2: Spoofing ICMP Packets . . . . .	2
3.3	Task 1.3: Traceroute . . . . .	3
3.4	Task 1.4: Sniffing and-then Spoofing . . . . .	4
<b>4</b>	<b>Task 2</b>	<b>6</b>

## List of Figures

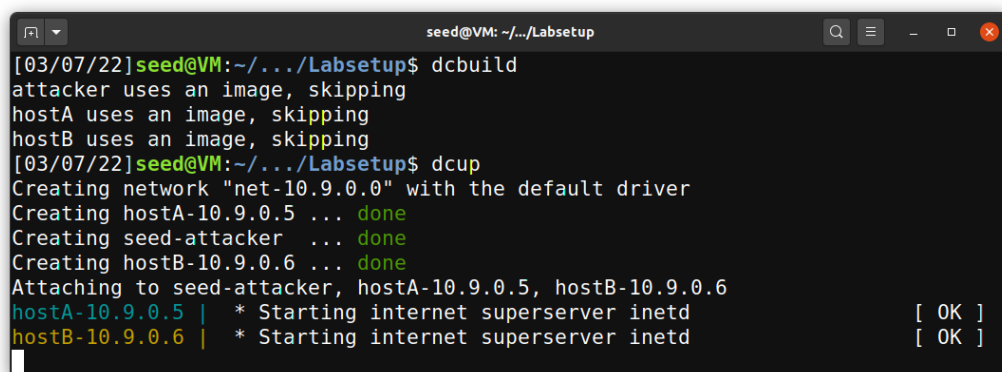
1	Setup of the Environment. . . . .	1
2	List of Docker network IDs. . . . .	1
3	Sniffing on the interface. . . . .	2
4	Sniffing on the interface with filters. . . . .	3
5	Spoofing ICMP Packets. . . . .	4
6	Traceroute with Scapy. . . . .	4
7	Examples of ping. . . . .	5
8	Examples of ping when running Sniffing and Spoofing. . . . .	6
9	Sniffer in C. . . . .	7

# 1 Introduction

Packet sniffing and spoofing are two important concepts in network security; they are two major threats in network communication. Being able to understand these two threats is essential for understanding security measures in networking. There are many packet sniffing and spoofing tools, such as Wireshark, Tcpdump, Netbox, Scapy, etc. Some of these tools are widely used by security experts, as well as by attackers. Being able to use these tools is important for students, but what is more important for students in a network security course is to understand how these tools work, i.e., how packet sniffing and spoofing are implemented in software.

## 2 Environment Setup

The environment, as suggested in the instructions, consists in a virtual machine running SEED-Ubuntu 20.04 on which are present three Docker containers, built and run using the `docker-compose.yml` file provided in the material.

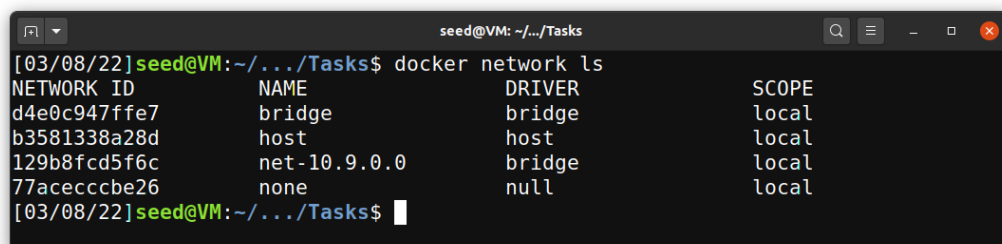


```
seed@VM: ~/.../Labsetup
[03/07/22]seed@VM:~/.../Labsetup$ dcbuild
attacker uses an image, skipping
hostA uses an image, skipping
hostB uses an image, skipping
[03/07/22]seed@VM:~/.../Labsetup$ dcup
Creating network "net-10.9.0.0" with the default driver
Creating hostA-10.9.0.5 ... done
Creating seed-attacker ... done
Creating hostB-10.9.0.6 ... done
Attaching to seed-attacker, hostA-10.9.0.5, hostB-10.9.0.6
hostA-10.9.0.5 | * Starting internet superserver inetd      [ OK ]
hostB-10.9.0.6 | * Starting internet superserver inetd      [ OK ]
```

Figure 1: Setup of the Environment.

### 2.1 Network Interface Name

After launching two shells for each of the host, we now want to retrieve the network interface name, action that has been done by running the `docker network ls` command on the virtual machine shell.



```
seed@VM: ~/.../Tasks
[03/08/22]seed@VM:~/.../Tasks$ docker network ls
NETWORK ID          NAME                DRIVER              SCOPE
d4e0c947ffe7        bridge             bridge              local
b3581338a28d        host               host                local
129b8fcd5f6c        net-10.9.0.0       bridge              local
77acecccbe26        none               null                local
[03/08/22]seed@VM:~/.../Tasks$
```

Figure 2: List of Docker network IDs.

As we can see, the name of the interface is `br-129b8fcd5f6c`.

## 3 Task 1

In this section, we will solve the given tasks regarding sniffing and spoofing.

### 3.1 Task 1.1: Using Scapy to Sniff and Spoof Packets

The code above will sniff the packets on the `br-129b8fcd5f6c` interface. For each captured packet, the callback function `print_pkt()` will be invoked; this function will print out some of the information about the packet.

```
1 #!/usr/bin/env python3
2
3 from scapy.all import *
4
5 def print_pkt(pkt):
6     print(' [PACKET SNIFFED] ')
7     print('      [SOURCE]: {} '.format(pkt[IP].src))
8     print('      [DESTINATION]: {} '.format(pkt[IP].dst))
9
10 print(' [SNIFFING...] ')
11 pkt = sniff(iface = 'br-129b8fcd5f6c', filter = 'icmp', prn = print_pkt)
```

We can now execute a ping command from the host A to the host B and see by running the sniffer on the virtual machine if we are able to see anything.

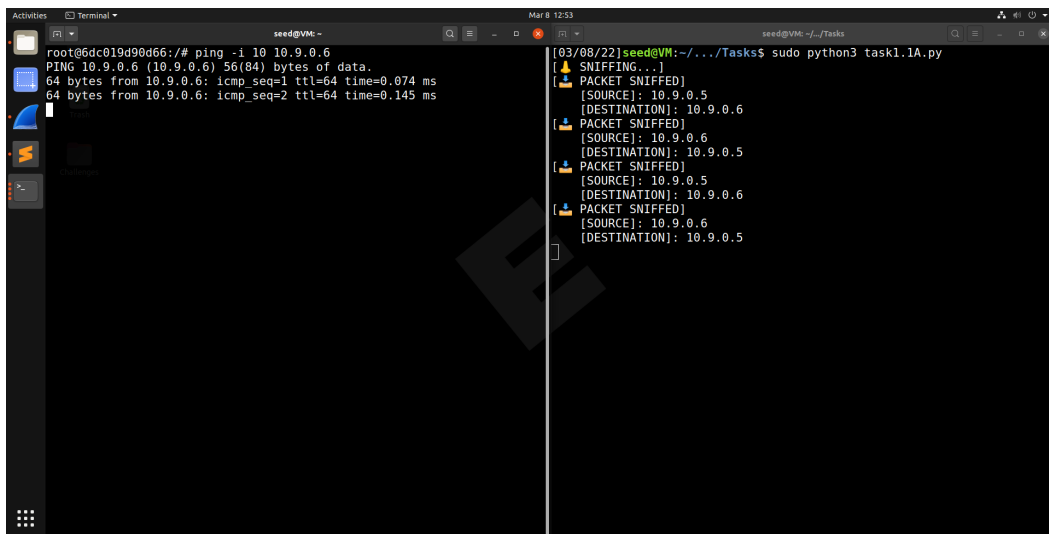


Figure 3: Sniffing on the interface.

As we can see, we are able to receive the packets. As a note, we used the `ping -i 10 10.9.0.6` just to be able to take a screenshot in time. Also, we are able to see both the ping and the reply.

Usually, when we sniff packets, we are only interested certain types of packets. We can do that by setting filters in sniffing. The code below allow to sniff only tcp packets on port 23 coming from the host with IP 10.9.0.5.

```
1 #!/usr/bin/env python3
2
3 from scapy.all import *
4
5 def print_pkt(pkt):
6     print(' [PACKET SNIFFED] ')
7     print('      [SOURCE]: {} '.format(pkt[IP].src))
8     print('      [DESTINATION]: {} '.format(pkt[IP].dst))
9
10 print(' [SNIFFING...] ')
11 pkt = sniff(iface='br-129b8fcd5f6c', filter='tcp port 23 and host 10.9.0.5', prn=print_pkt)
```

The results are shown in Fig. (4). As we can notice, we are able to see traffic, while with ping commands we wouldn't be able to see anything.

### 3.2 Task 1.2: Spoofing ICMP Packets

As a packet spoofing tool, Scapy allows us to set the fields of IP packets to arbitrary values. The objective of this task is to spoof IP packets with an arbitrary source IP address. We will spoof ICMP echo request packets,



Figure 4: Sniffing on the interface with filters.

and send them to another VM on the same network. We will use Wireshark to observe whether our request will be accepted by the receiver. If it is accepted, an echo reply packet will be sent to the spoofed IP address. On the virtual machine, we are executing the following code:

```
1 #!/usr/bin/env python3
2
3 from scapy.all import *
4
5 ip = IP()
6
7 src = '10.9.0.5'
8 dst = '10.9.0.6'
9
10 ip.src = src
11 ip.dst = dst
12
13 icmp = ip/ICMP()
14
15 print('[SENDING PACKET]')
16 print('    [SOURCE]: {}'.format(src))
17 print('    [DESTINATION]: {}'.format(dst))
18
19 send(icmp, verbose = 0)
20 print('[PACKET SENT]')
```

Since on Wireshark we are listening on the br-129b8fcd5f6c interface, we are using the IPs of the containers connected to it. Results are found in Fig. (5).

### 3.3 Task 1.3: Traceroute

The objective of this task is to use Scapy to estimate the distance, in terms of number of routers, between our VM and a selected destination. This is basically what is implemented by the traceroute tool. In this task, we will write our own tool. The idea is quite straightforward: just send a packet (any type) to the destination, with its Time-To-Live (TTL) field set to 1 first. This packet will be dropped by the first router, which will send us an ICMP error message, telling us that the time-to-live has exceeded. We can do this automatically using a simple loop in Python, as in the following code.

```
1 #!/usr/bin/env python3
2
3 from scapy.all import *
4
5 dst = 'google.com'
6
7 # Starting from 1
8 ttl = 1
9
```

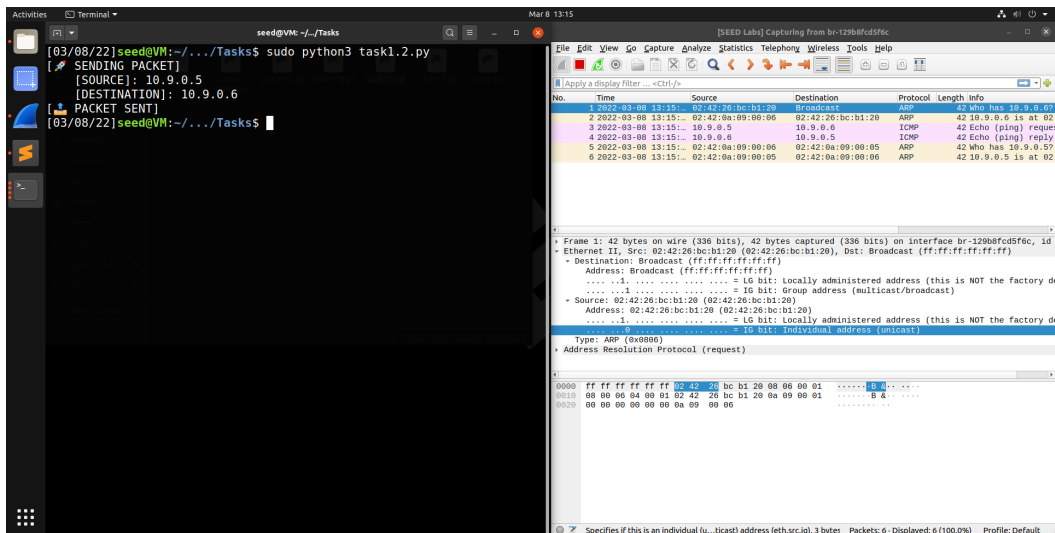


Figure 5: Spoofing ICMP Packets.

```

10 while 1:
11     pkt = sr1(IP(dst = dst, ttl = ttl)/ICMP(), verbose = 0) # Verbose handle the output
12     # If the TTL is exceeded
13     if pkt[ICMP].type == 11 and pkt[ICMP].code == 0:
14         print('TTL = {} IP = {}'.format(ttl, pkt.src))
15         ttl += 1
16     # Otherwise we reached our destination
17     elif pkt[ICMP].type == 0:
18         print('REACHED W/ TTL = {} IP = {}'.format(ttl, pkt.src))
19         break

```

The output is shown in Fig. (6) and as we can see we are able to determine the IP address of every machine inside the route between our virtual machine and the destination google.com, as well as the number of routers between them (in this case 10).

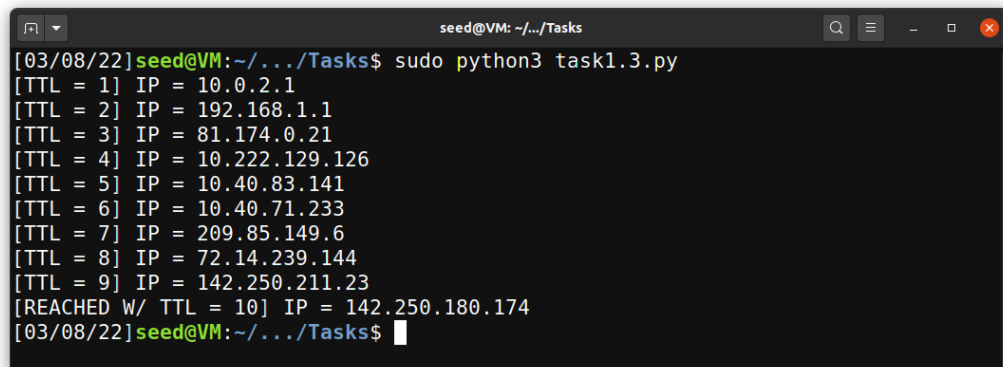


Figure 6: Traceroute with Scapy.

### 3.4 Task 1.4: Sniffing and-then Spoofing

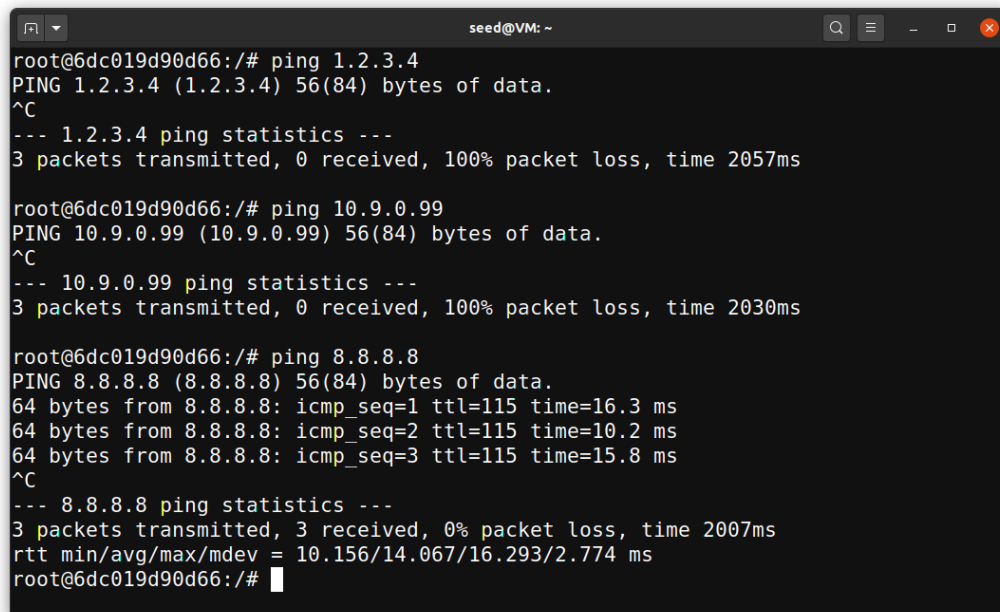
In this task, we will combine the sniffing and spoofing techniques to implement the following sniff and then spoof program. From the user container, we ping an IP X. This will generate an ICMP echo request packet. If X is alive, the ping program will receive an echo reply, and print out the response. Our sniff-and-then-spoof program runs on the VM, which monitors the LAN through packet sniffing. Whenever it sees an ICMP echo request, regardless of what the target IP address is, our program should immediately send out an echo reply using the packet spoofing technique. Therefore, regardless of whether machine X is alive or not, the ping program will always receive a reply, indicating that X is alive. The code used is the following:

```

1 #!/usr/bin/env python3
2
3 from scapy.all import *
4
5 def spoof(pkt):
6     """
7     Given a received packet, spoof its address and send a reply
8     """
9
10    # Execute only if the packet is an echo request
11    if pkt[ICMP].type != 8:
12        return
13    else:
14        # Destination becomes source and viceversa
15        spoof_ip = IP(src = pkt[IP].dst, dst = pkt[IP].src)
16
17        spoof_icmp = ICMP(type = 0, id = pkt[ICMP].id, seq = pkt[ICMP].seq)
18
19        spoof_data = pkt[Raw].load
20
21        spoof_pkt = spoof_ip/spoof_icmp/spoof_data
22        send(spoof_pkt, verbose = 0)
23
24        print(' [SPOOFED PACKET SENT] ')
25        return
26
27 print(' [SNIFFING...] ')
28 pkt = sniff(iface = 'br-129b8fcd5f6c', filter = 'icmp', prn = spoof)

```

We had to add the first `if pkt[ICMP].type != 8:` otherwise it would have spoofed also echo replies (if the type of ICMP is 8 then it's an echo request, else if it's 0 then it's an echo reply). When we spoof we IP the invert the source and the destination (since it's a reply) and we set the ICMP type to 0 (since it would be an echo reply). We also keep everything else untouched as the ICMP id and seq, and we further attach the data by calling `pkt[Raw].load`. First, let's see what happens when pinging a non-existing host on the Internet (1.2.3.4), a non-existing host on the LAN (10.9.0.99) and an existing host on the Internet (8.8.8.8).



```

seed@VM: ~
root@6dc019d90d66:/# ping 1.2.3.4
PING 1.2.3.4 (1.2.3.4) 56(84) bytes of data.
^C
--- 1.2.3.4 ping statistics ---
3 packets transmitted, 0 received, 100% packet loss, time 2057ms

root@6dc019d90d66:/# ping 10.9.0.99
PING 10.9.0.99 (10.9.0.99) 56(84) bytes of data.
^C
--- 10.9.0.99 ping statistics ---
3 packets transmitted, 0 received, 100% packet loss, time 2030ms

root@6dc019d90d66:/# ping 8.8.8.8
PING 8.8.8.8 (8.8.8.8) 56(84) bytes of data.
64 bytes from 8.8.8.8: icmp_seq=1 ttl=115 time=16.3 ms
64 bytes from 8.8.8.8: icmp_seq=2 ttl=115 time=10.2 ms
64 bytes from 8.8.8.8: icmp_seq=3 ttl=115 time=15.8 ms
^C
--- 8.8.8.8 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2007ms
rtt min/avg/max/mdev = 10.156/14.067/16.293/2.774 ms
root@6dc019d90d66:/#

```

Figure 7: Examples of ping.

As we can expect, there are no packets received when we try to ping non existent IP (both on the internet or on the LAN), while we have responses when pinging an existing address. Let's now run our sniff-and-spoof program and see the output, which can be found in Fig. (8). We can see that when pinging a non-existing host on the Internet (1.2.3.4), our program is able to spoof the address and send replies which are correctly received by the host. When pinging a non-existing host on the LAN (10.9.0.99) however we have

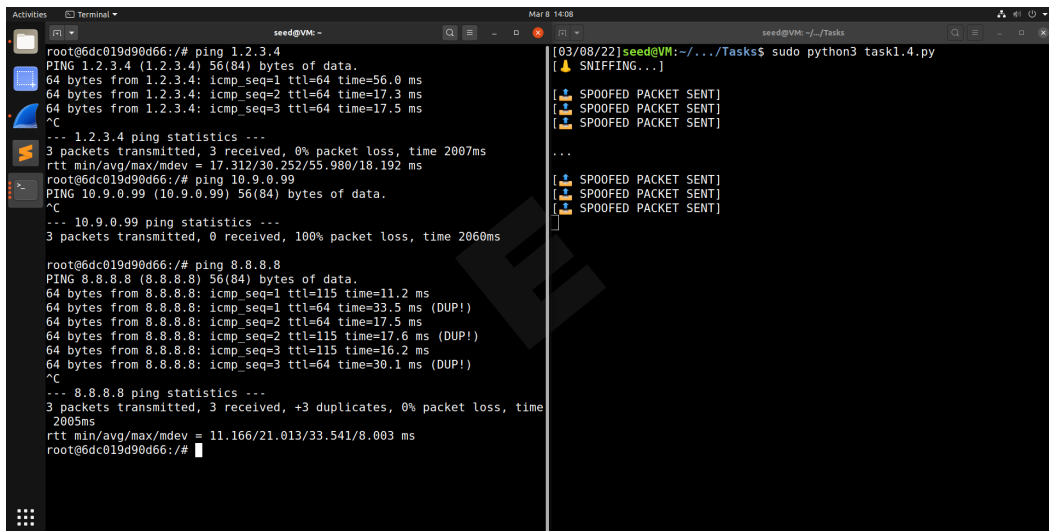


Figure 8: Examples of ping when running Sniffing and Spoofing.

no response: this happens because no ARP cache entry exists for a requested destination IP, and therefore the kernel will generate ARP requests until receiving an answer, thus failing our spoofing attack. Finally, when pinging an existing host on the Internet (8.8.8.8), we can see that the host receives a duplicate of the reply (marked with DUP!), since it's getting a reply both from the legitimate source and from our program.

## 4 Task 2

Even though I've little experience with the C language, I tried to complete the second task nonetheless. Using the pcap library, I tried to implement a sniffer that print out the source and destination IP addresses of each captured packet. This has been done by modifying the provided code in the following way.

```

1 #include <pcap.h>
2 #include <stdio.h>
3 #include <stdlib.h>
4
5 // Ethernet Header
6 struct ethheader {
7     u_char ether_dhost[6]; // Destination address
8     u_char ether_shost[6]; // Source address
9     u_short ether_type;    // Type
10 };
11
12 // IP Header
13 struct ipheader {
14     unsigned char iph_ihl:4, // IP header length
15                 iph_ver:4;   // IP version
16     unsigned char iph_tos;    // Type of service
17     unsigned short int iph_len; // IP Packet length (data + header)
18     unsigned short int iph_ident; // Identification
19     unsigned short int iph_flag:3, // Fragmentation flags
20                 iph_offset:13; // Flags offset
21     unsigned char iph_ttl;    // Time to Live
22     unsigned char iph_protocol; // Protocol type
23     unsigned short int iph_checksum; // IP datagram checksum
24     struct in_addr iph_sourceip; // Source IP address
25     struct in_addr iph_destip;  // Destination IP address
26 };
27
28 /* This function will be invoked by pcap for each captured packet.
29 We can process each packet inside the function. */
30 void got_packet(u_char *args, const struct pcap_pkthdr *header,
31 const u_char *packet) {
32     printf("[PACKET RECEIVED]\n");
33     // Get the header information of Ethernet and IP
34     struct ethheader *eth = (struct ethheader *)packet;

```



```

35 struct ipheader *ip_pkt = (struct ipheader *) (packet + sizeof(struct ethheader));
36 printf("    [SOURCE IP]: %s \n", inet_ntoa(ip_pkt->iph_sourceip));
37 printf("    [SOURCE IP]: %s \n", inet_ntoa(ip_pkt->iph_destip));
38 }
39
40 int main() {
41     pcap_t *handle;
42     char errbuf[PCAP_ERRBUF_SIZE];
43     struct bpf_program fp;
44     char filter_exp[] = "icmp";
45     bpf_u_int32 net;
46
47     // Step 1: Open live pcap session on NIC with name br-129b8fcd5f6c.
48     handle = pcap_open_live("br-129b8fcd5f6c", BUFSIZ, 1, 1000, errbuf);
49
50     // Step 2: Compile filter_exp into BPF psuedo-code
51     pcap_compile(handle, &fp, filter_exp, 0, net);
52
53     if (pcap_setfilter(handle, &fp) != 0) {
54         pcap_perror(handle, "[ERROR]: ");
55         exit(EXIT_FAILURE);
56     }
57
58     // Step 3: Capture packets
59     pcap_loop(handle, -1, got_packet, NULL);
60     pcap_close(handle);
61     return 0;
62     //Close the handle
63 }
64 // Note: don't forget to add "-lpcap" to the compilation command.
65 // For example: gcc -o sniff sniff.c -lpcap

```

The main contribution of this code is the addition of the first two struct, that define the different elements of the Ethernet and the IP header (even though we are only interested in the IP source address and the IP destination address). The proof of work is given in Fig. (9).

The screenshot shows a terminal window with two panes. The left pane shows the output of a ping command: `root@6dc019d90d66:/# ping -i 10 10.9.0.6`, followed by three successful ping responses from 10.9.0.6. The right pane shows the output of a sniffer program: `[03/10/22]seed@VM:~/.../Tasks$ sudo ./sniffer`, followed by a series of "PACKET RECEIVED" messages, each showing the source IP as either 10.9.0.5 or 10.9.0.6.

Figure 9: Sniffer in C.